

Test-suite for Automating Data-consistency checks on HBase

Pradeep S, Mallikarjun V
Flipkart

About Flipkart

- India's largest online retailer:
 - 10M page visits a day
 - 2M shipments a day
 - 30M products across more than 70 categories
- Big Billion Days (\$300M sales, top ranked app on Google Play Store)

Agenda

- Yak: HBase Cluster @ Flipkart
 - Need for scalable, multi-tenant, strongly consistent data stores
- Yak: Need for Data-correctness Guarantees
- Why Data-consistency Test-suite?
- How we did the Test-suite?

Yak: HBase Cluster @ Flipkart

- HBase for OLTP Key-Value store
- Bring-Your-Own-Box multi-tenancy on HBase
- RSGroup based isolation on HBase 1.2.4 & 2.1.3
- WAL based Change-Data-Capture into Kafka
- Stores critical data-sets in e-commerce like: Orders, Payments

Yak: Data-correctness Guarantees Required

- Read-Your-Own-Write consistency at single row level
- 'Atleast-once' guarantee in the change-data-capture
- Ordering guarantee in the change-data-capture
- Predictable recovery times, no data-loss upon failures
- Data-reliability for HBase admin operations

Why Data-consistency Test-suite?

- To enable faster upgrade cycles
- Releases don't degrading on the data-guarantees
- To assist in reproducing edge-cases scenarios
- To accommodate additional failure scenarios

Different Approaches for the Test-suite

- Algorithm testing: TLA+, etc.
- Code testing: Test suite of HBase, Hadoop, Zookeeper
- Testing on running cluster: Jepsen, ChaosMonkey & ITBLL of HBase, etc.

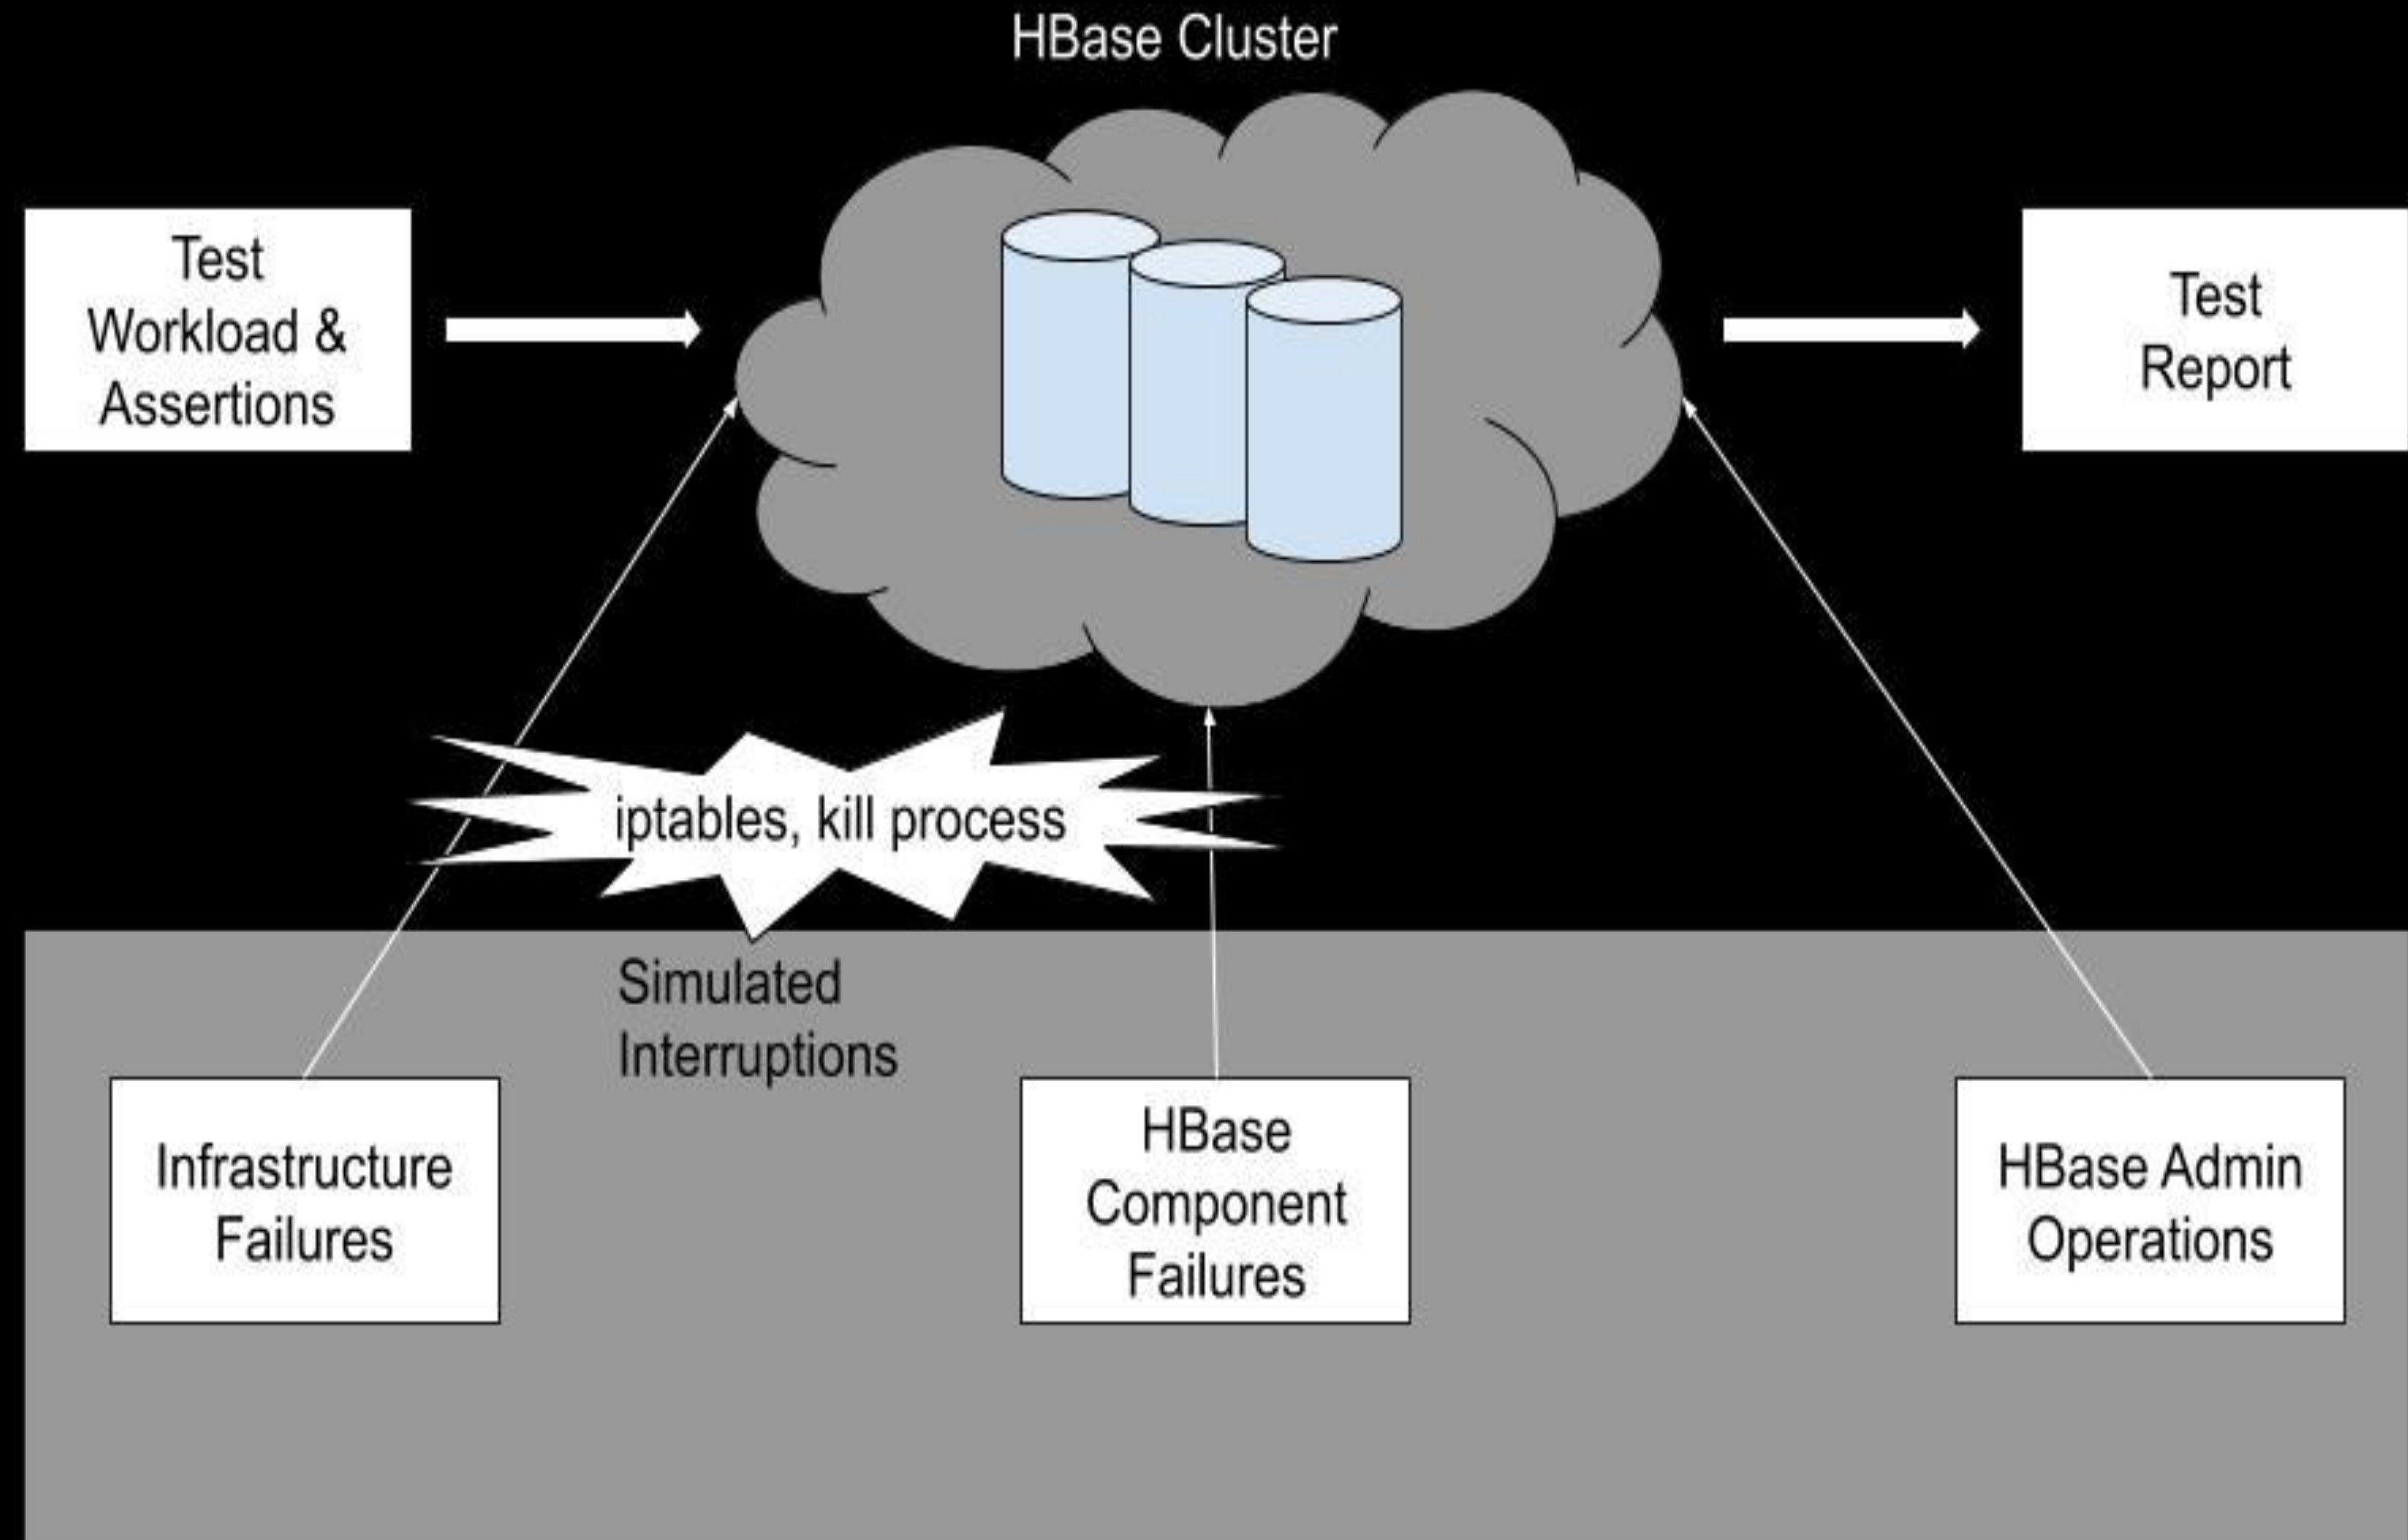
Different Approaches for the Test-suite

- ~~Algorithm testing: TLA+, etc.~~ - Implementation, Deployment NOT tested
- ~~Code testing: Test suite of HBase~~ - Deployment, Integration NOT tested
- Testing on running cluster: Jepsen, ChaosMonkey & ITBLL of HBase, etc.

Reason: Integrated test with { Algorithm + Code + Deployment }

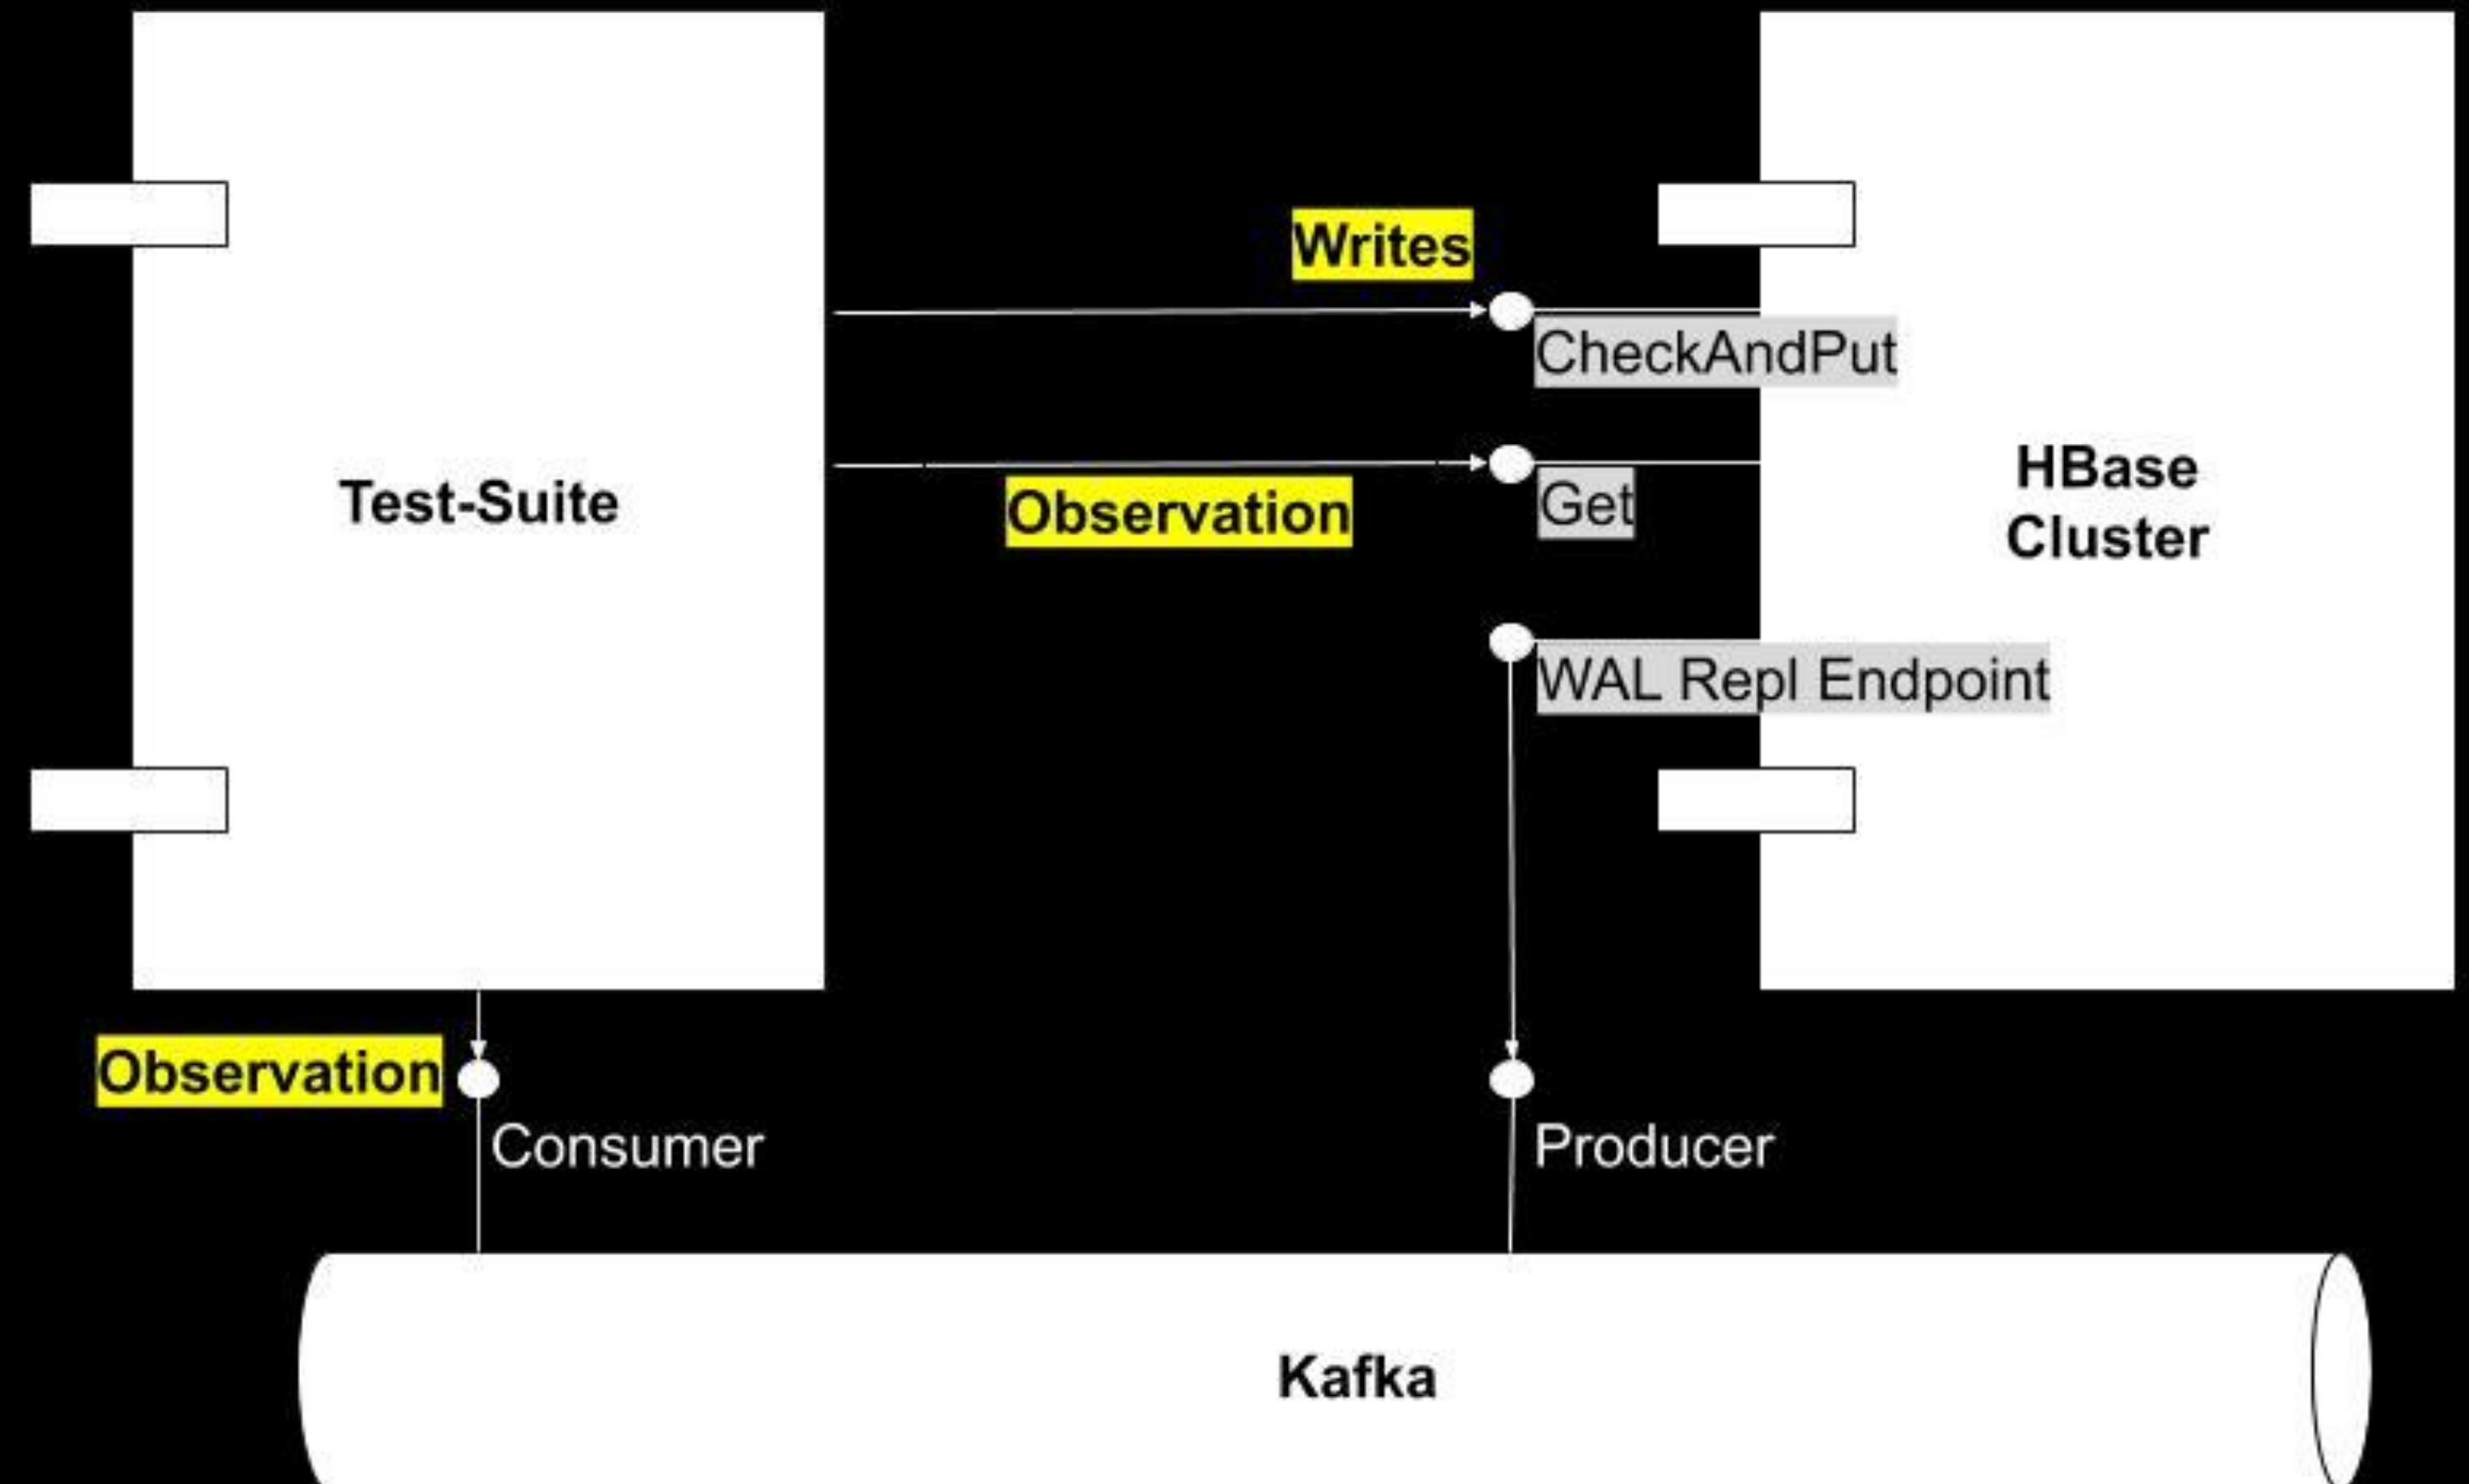
Test-Suite

- Jepsen, ChaosMonkey based
- Test workload on HBase
- Simulated Interrupts:
 - Infrastructure failures
 - HBase component failures
 - Admin operations
- Test Report

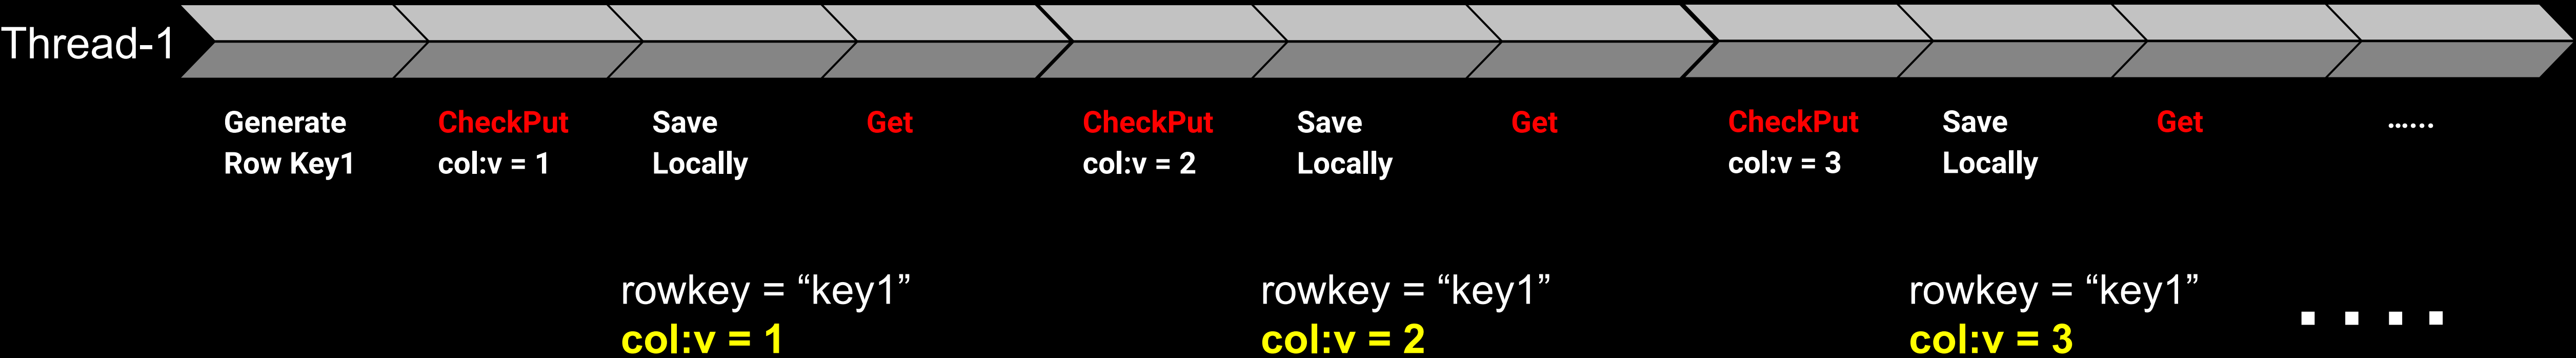


Test-Suite - Test Workload

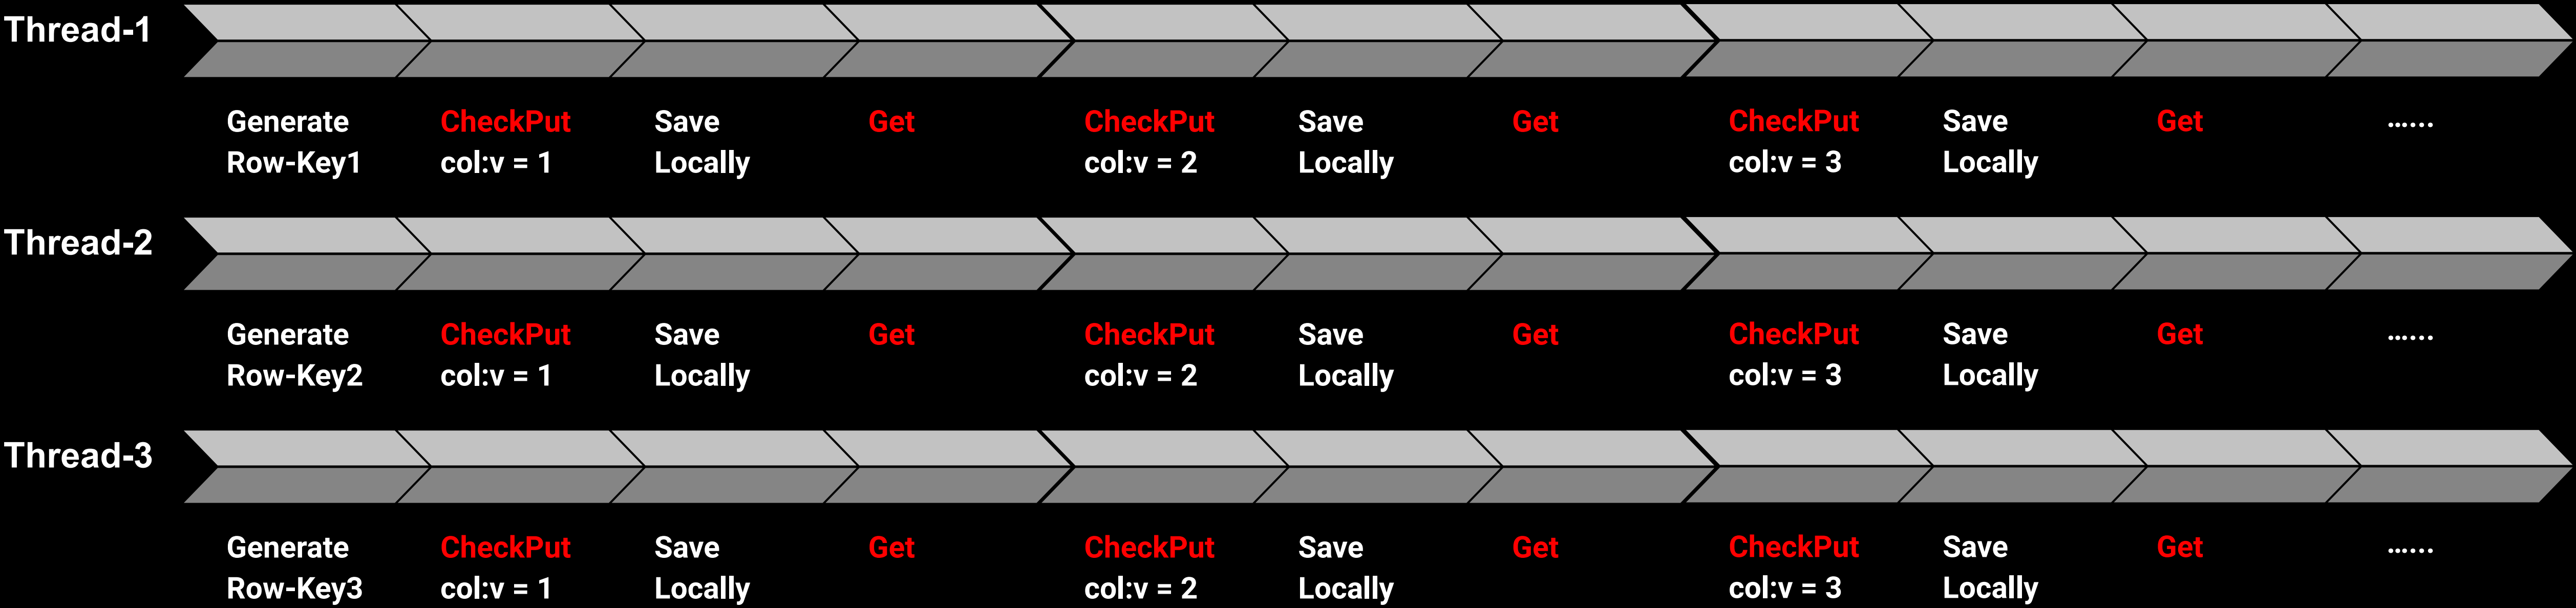
- Writes using CheckAndPut
- Observation points:
 - HBase Get
 - *Kafka consumer from CDC*



Test-Suite - Test Workload




Test-Suite - Test Workload



N Threads... on M Machines

Test-Suite - Assertions

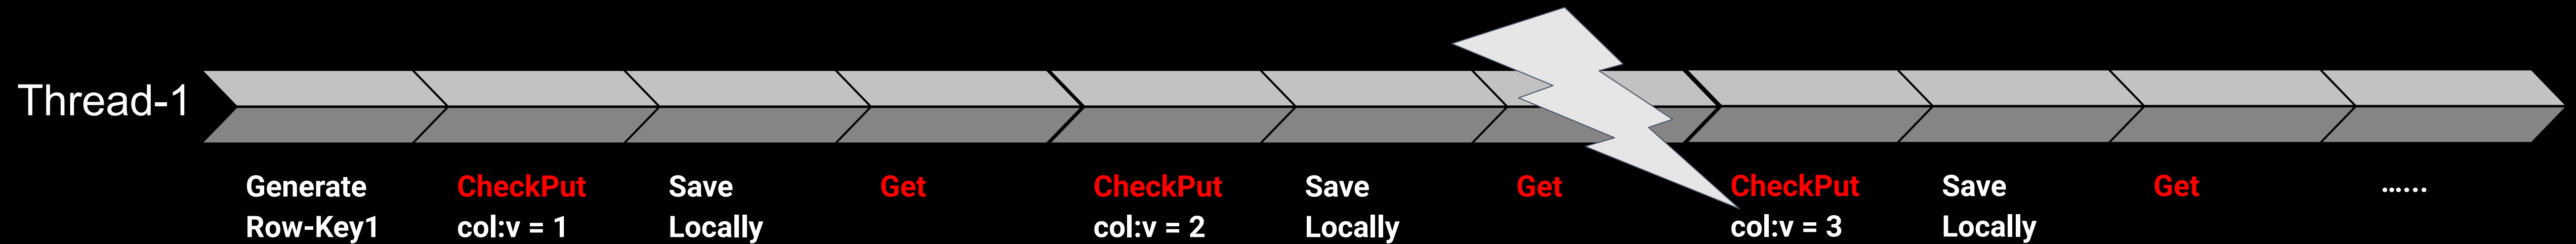


- Locally saved-state is compared against Get call result for no data-loss

For Key1 after every **CheckAndPut**:

```
Assert ( ThreadLocal<Version> == Version from HBase )
```

Test-Suite - Assertions



- Kafka listener reads the events of the key:

For a key1 having writes as: 1,2,3,4

1,2,2,2,3,4

1,2,3,2,3,4


1,2,4,2,3,4 **×**

1,2,4 **×**

For the key1:

```
Assert (
  read event version <=
  latest seen version +1
)
```

Test-Suite - Assertions




- HBase regions of a table are within the RSGroup nodes
- HDFS data-blocks stored within the RSGroup nodes

Assert (Region Assignment within RSGroup Nodes)

Assert (HDFS blocks within RSGroup Nodes)

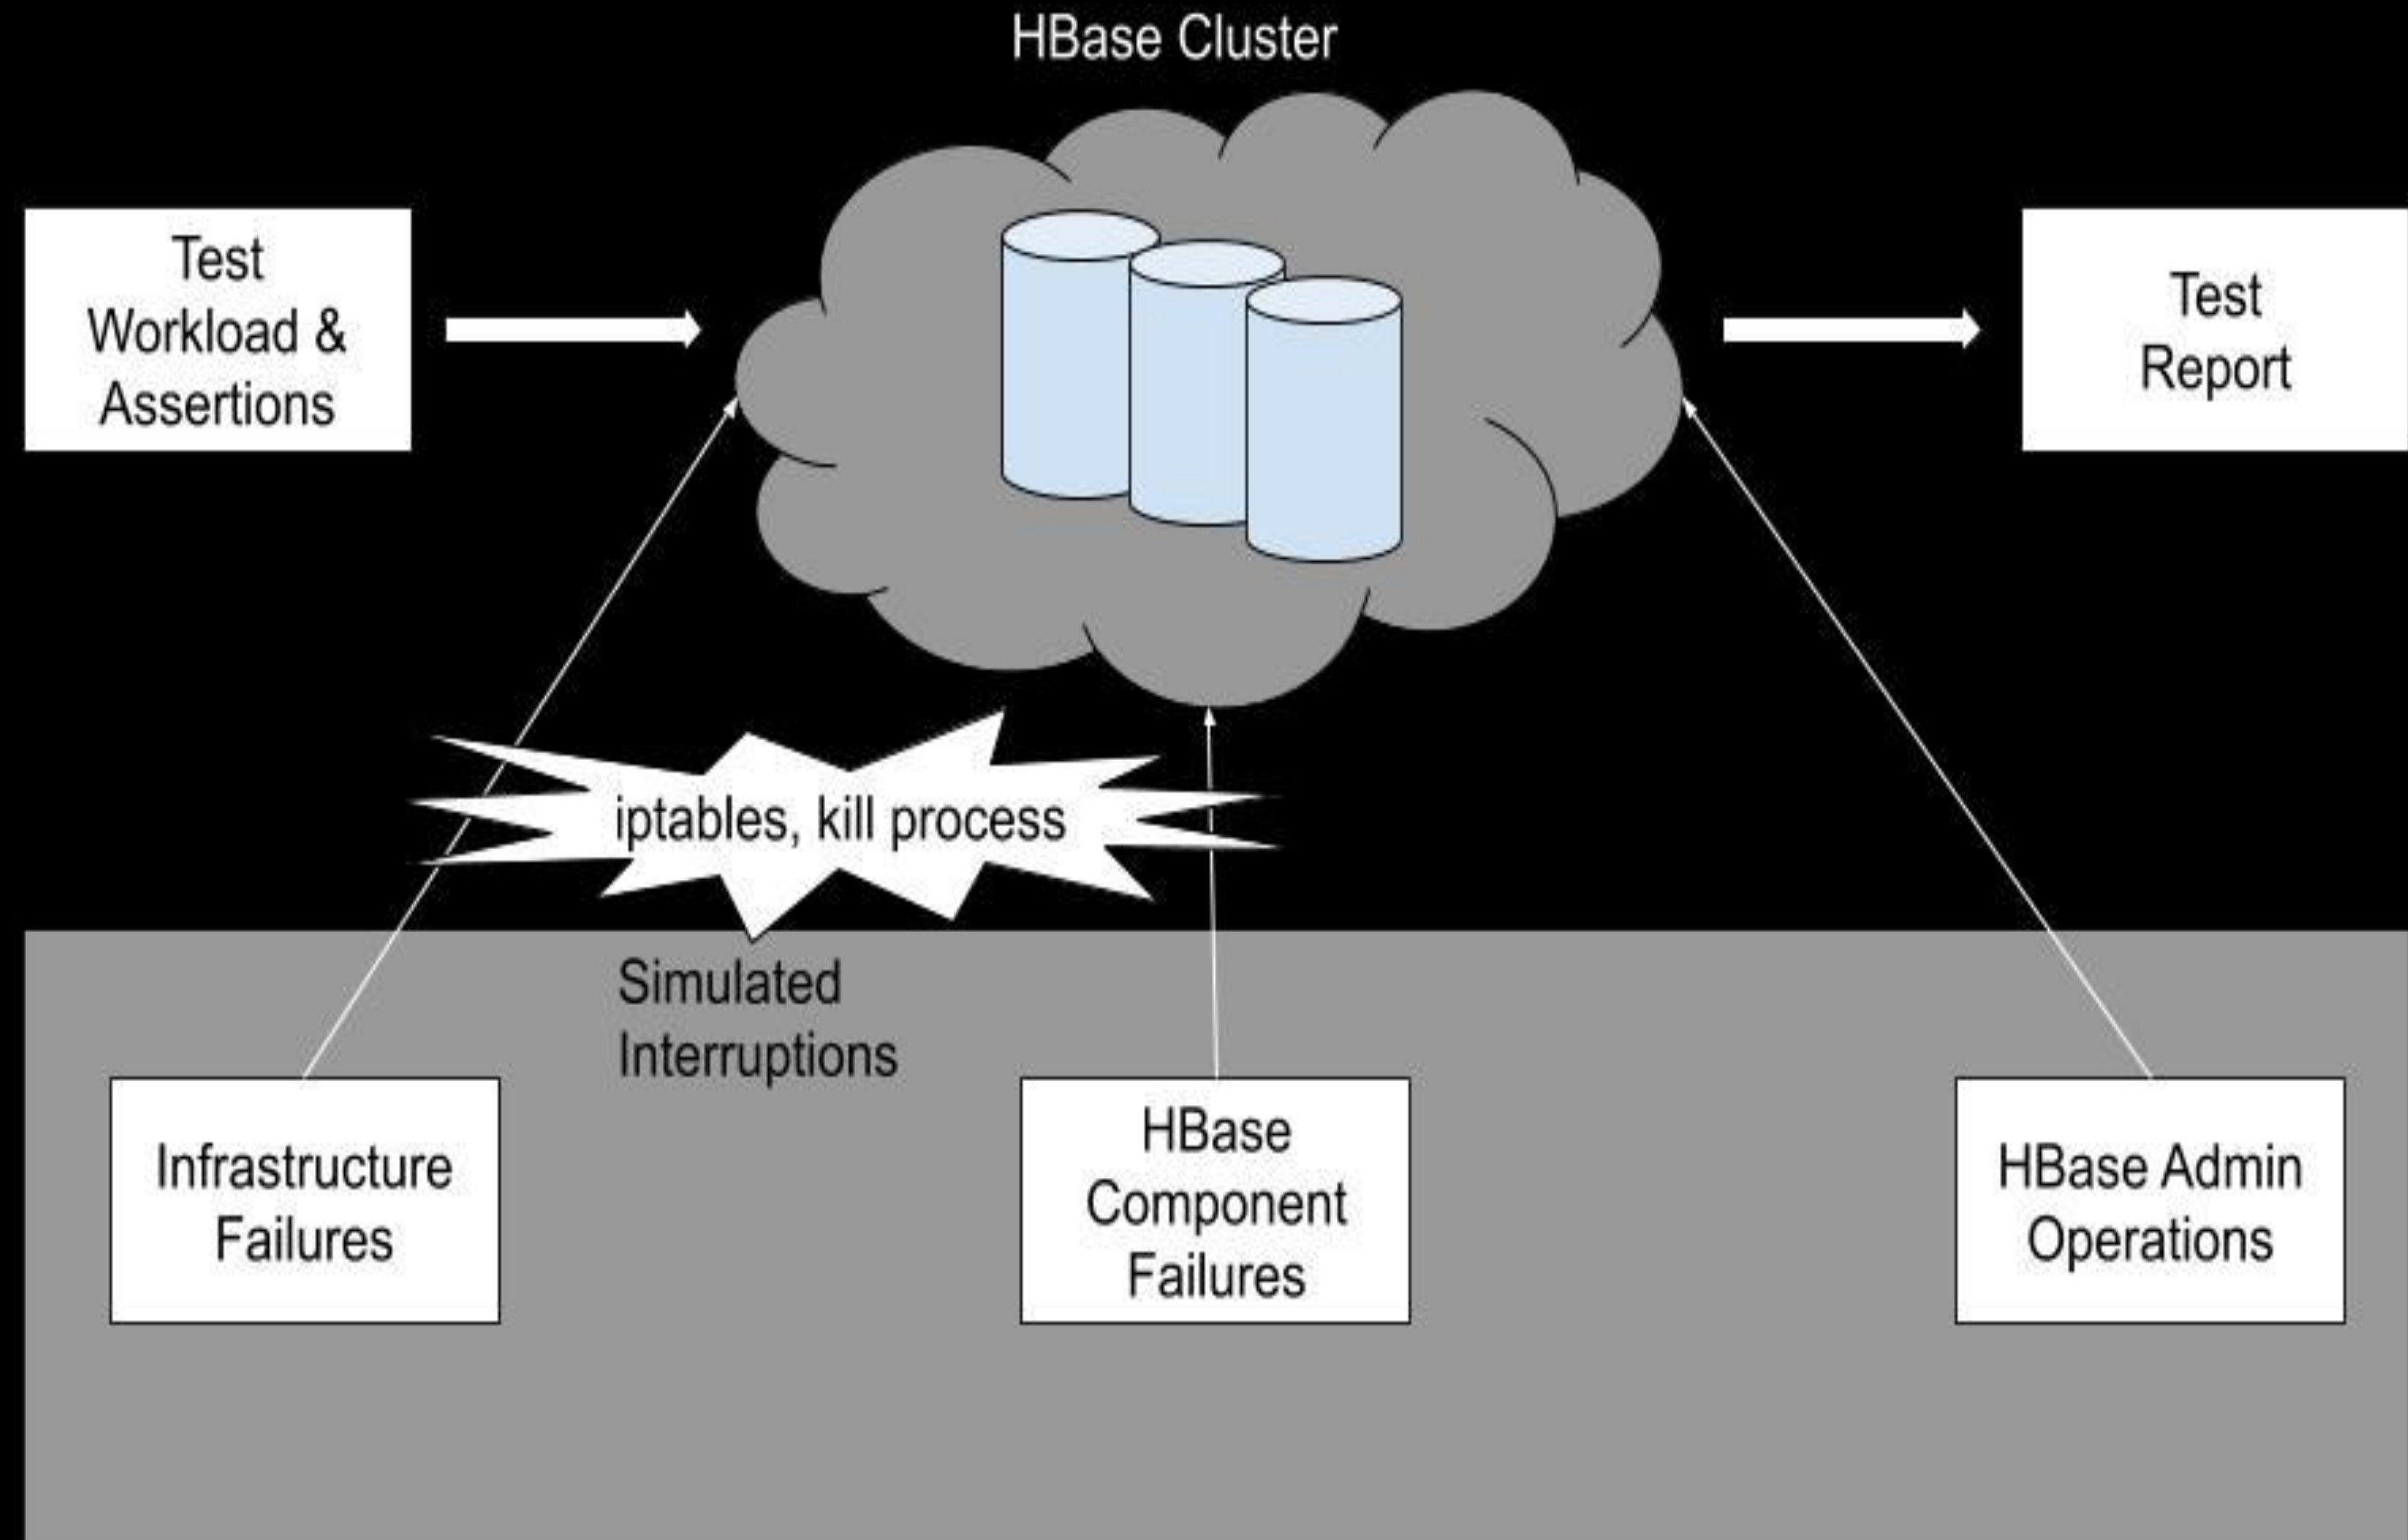
Test-Suite - Assertions



- Locally saved-state is compared against Get call result for no data-loss
- Kafka listener reads the events of the key:
 - To assert no data-loss as the versions are incremental
 - To assert no-ordering loss
- RSGroup isolation checks: WAL & HDFS data-blocks & Region Assignment

Interruptions

- Simulate issue & heal
- Types of Interruptions:
 - Infrastructure failures
 - HBase component failures
 - HBase admin operations



Interruptions - Infrastructure Failures

- Network failures:
 - Network partition/failure: within zookeeper, region-servers, namenode etc.
 - Packet loss: within zookeeper, master, region-servers, namenode etc.
 - Tools: iptables, tc, comcast
- Other failure modes - to be added:
 - Clock skew
 - Packet delays
 - Disk read-only etc

Interruptions - Component Failures

- **Kill hbase components:**
 - region-server, master-node, zookeeper
- **Kill hadoop components:**
 - data-node, journal-node, name-node
- **Node crash:**
 - region-server-node, master-node, name-node, zookeeper, journal-node

Interruptions - HBase Admin Operations

- Split/merge region
- Assign/move region
- Restart: region-server, data-node
- Stop: region-server, data-node, etc.
- Kafka properties reload for CDC

Issues Uncovered & Fixed

- Assign an already assigned region causing data-loss - Fixed in 2.x
- Ordering loss in tail process from WAL file upon a region failover - Fixed in 2.x
- WAL not-isolated across region-server group - [HBASE-21641](#)

Further Plans

- Migration to HBase ChaosMonkey for Interruptions
- Opensource the test-suite
- WAL push to Kafka - <https://github.com/flipkart-incubator/hbase-sep>

Thanks !

Appendix

Limitations

- Tests are not deterministic. Needs 100s of iterations.
- Doesn't catch all the bugs. But catches the practical issues for the specific use-case.

Test-Suite - Sample Report

Yak Testbed (TestCount - 58)																						
Test Name		Interruptions				HbaseMetrics Output			KafkaMetrics Output			Running Time				Process Status						Overall Status
PutGetKafkaTest	mergeRegions	SUCCESS	3	Durations			Metric	Status	Failures	Metric	Status	Failures	RunningTime	StartTime	EndTime	Duration(ms)	PASSED					
				dataMismatchRH	YES	98	dataMismatchRH	YES	98													
				dataMismatchWH	NO	0	dataMismatchWH	NO	0													
				dataLoss	NO	0	dataLoss	NO	0													
	assignRegions	SUCCESS	2	StartTime	EndTime	Duration(ms)	checkPutFail	YES	264	connectionMismatch	NO	0	getTime	08:35:44	08:37:17	92813						
				08:36:02	08:36:30	27536	checkPutException	NO	0	dataMismatchWH	NO	0	putTime	08:35:44	08:37:16	92077						
				08:36:02	08:36:24	21678	getFail	NO	0	dataLoss	NO	0	kafkaTime	08:35:44	08:37:55	131393						
				inconsistencies	NO	0	ordering	YES	0													
				wallisolated	YES	0	repetition	NO	0													
				hBlockIsolated	NO	198																
allRegionsOpen	YES	0																				
failuresAfterTest	NO	0																				
PutGetKafkaTest	mergeRegions	SUCCESS	3	Durations			Metric	Status	Failures	Metric	Status	Failures	RunningTime	StartTime	EndTime	Duration(ms)	PASSED					
				dataMismatchRH	YES	98	dataMismatchRH	YES	98													
				dataMismatchWH	NO	0	dataMismatchWH	NO	0													
				dataLoss	NO	0	dataLoss	NO	0													
	assignRegions	SUCCESS	2	StartTime	EndTime	Duration(ms)	checkPutFail	YES	264	connectionMismatch	NO	0	getTime	08:35:44	08:37:17	92813						
				08:36:02	08:36:30	27536	checkPutException	NO	0	dataMismatchWH	NO	0	putTime	08:35:44	08:37:16	92077						
				08:36:02	08:36:24	21678	getFail	NO	0	dataLoss	NO	0	kafkaTime	08:35:44	08:37:55	131393						
				inconsistencies	NO	0	ordering	YES	0													
				wallisolated	YES	0	repetition	NO	0													
				hBlockIsolated	NO	198																
allRegionsOpen	YES	0																				
failuresAfterTest	NO	0																				
PutGetKafkaTest	mergeRegions	SUCCESS	3	Durations			Metric	Status	Failures	Metric	Status	Failures	RunningTime	StartTime	EndTime	Duration(ms)	PASSED					
				dataMismatchRH	YES	98	dataMismatchRH	YES	98													
				dataMismatchWH	NO	0	dataMismatchWH	NO	0													
				dataLoss	NO	0	dataLoss	NO	0													
	assignRegions	SUCCESS	2	StartTime	EndTime	Duration(ms)	checkPutFail	YES	264	connectionMismatch	NO	0	getTime	08:35:44	08:37:17	92813						
				08:36:02	08:36:30	27536	checkPutException	NO	0	dataMismatchWH	NO	0	putTime	08:35:44	08:37:16	92077						
				08:36:02	08:36:24	21678	getFail	NO	0	dataLoss	NO	0	kafkaTime	08:35:44	08:37:55	131393						
				inconsistencies	NO	0	ordering	YES	0													
				wallisolated	YES	0	repetition	NO	0													
				hBlockIsolated	NO	198																
allRegionsOpen	YES	0																				
failuresAfterTest	NO	0																				